

# Processes

processes practice for CSC420 Operating Systems Spring 2024 Lyon  
College

Marcus Birkenkrahe

April 23, 2024

## README

- This file accompanies lectures on the shell and `bash(1)`. To gain practice, you should type along in your own Org-mode file. You have to have Emacs and my `.emacs` file installed on your PC or the Pi you're working with.
- This section is based on chapter 10 of Shotts, The Linux Command Line (2e), NoStarch Press (2019).
- To make this easier, use the auto expansion (`<s`). This will only work if you have my `.emacs` file (from GDrive) installed.
- Add the following two lines at the top of your file, and activate each line with `C-c C-c` (this is confirmed in the echo area as `Local setup has been refreshed`):

```
#+PROPERTY: header-args: bash :results output
```

- Remember that `C-M-\` inside a code block indents syntactically (on Windows, this may only work if you have a marked region - set the mark with `C-SPC`).

## Overview

- Modern operating systems are *multitasking*, which means they create an illusion of doing more than one thing at once.

- They do this by rapidly switching from one executing program to another.
- The **kernel** manages this through clever **process management**, which really is clever **memory management**.
- This is illustrated in the figure 1. The simple program is all over the computer's memory.

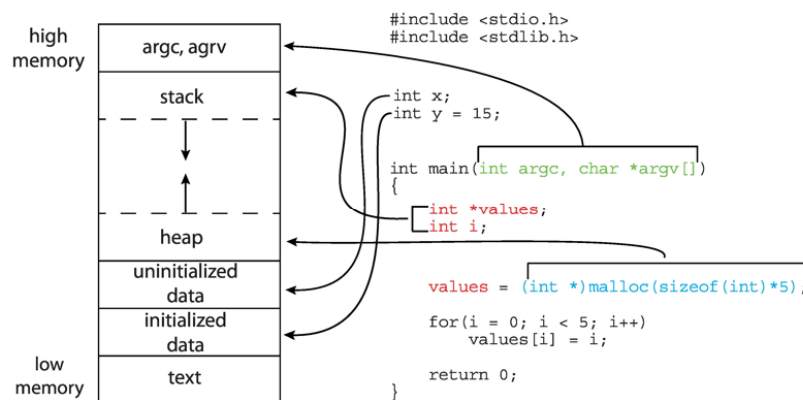


Figure 1: Memory Layout of a C Program (Source: Silberschatz et al)

- The diagram illustrates the typical memory layout of a process in a C program:
  1. **High Memory to Low Memory:** The top represents high memory addresses and the bottom low memory addresses in the process's address space.
  2. **Text:** Contains the compiled code of the program (`.out`).
  3. **Uninitialized Data:** Holds uninitialized global and static variables (values unknown at start) - `int x;`
  4. **Initialized Data:** Contains initialized global and static variables - `int y = 15;`
  5. **argc, argv:** Located on the stack; used to pass command-line arguments into the program. `argc` is the argument count, and `*argv[]` is an pointer array of argument strings.

6. **Stack:** Used for function call management; contains local variables and function parameters - `int *values` (pointer to `values`), `int i`.
7. **Heap:** For dynamically-allocated memory during runtime - `malloc` allocates memory for 5 integers, assigned to `values`.

x

- As always, let's focus on stuff we can do ourselves. This includes a bunch of new shell commands:

COMMAND	MEANING
<code>ps</code>	Report a snapshot of current processes
<code>top</code>	Display tasks
<code>jobs</code>	List active jobs
<code>bg</code>	Place a job in the background
<code>fg</code>	Place a job in the foreground
<code>kill</code>	Send a kill signal to process
<code>killall</code>	Kill processes by name
<code>shutdown</code>	Shut down or reboot the system

## How a process works

- When the OS starts up, the **kernel** launches the `init` program, which in turn runs a series of shell scripts (in `/etc`) that start all the system services.
- Check `/etc` out now - you find e.g. the directory `/etc/cups`, which contains scripts for the Common UNIX Printing System (CUPS).
- Many of the services are *daemon* programs - they just sit in the background and do their thing without a user interface (UI).
- `init` itself is a **daemon**, also called `systemd`. The shell program `systemctl` allows indirect access to all services.

- **Grab a daemon!**

In the code block ??,

1. run the command `systemctl status`,
2. tee its output to a text file `systemctl.txt`

### 3. `grep` for the login daemon program `logind`

```
systemctl status | tee systemctl.txt | grep logind
```

```
      86093 grep logind
systemd-logind.service
      874 /lib/systemd/systemd-logind
```

- If a program (like `init`) can launch other programs, it's a *parent process* producing a *child process*.
- **How does the kernel maintain control?** By assigning a *process ID* (PID) to every process.
- Processes are assigned in ascending order beginning with `init`, which has PID 1.
- Processes are assigned in ascending order beginning with `init`, which has PID 1: run `ps ax`, `grep` for `init`, and print the first line:

```
ps ax | grep init | head -n 1
```

```
1 ?          Ss      0:05 /sbin/init splash
```

- The **kernel** also tracks process memory and readiness to resume execution. Like files, processes have owners and userIDs.

## Viewing processes statically

- The `ps` program has a lot of options (check `ps(1)`)
- Run `ps` without options.

```
ps | head -n 10
```

PID	TTY	TIME	CMD
1313	?	00:00:00	systemd
1314	?	00:00:00	(sd-pam)
1323	?	00:00:00	pipewire
1324	?	00:47:15	pulseaudio

```

1325 ?      00:00:01 cinnamon-sessio
1343 ?      00:00:04 dbus-daemon
1541 ?      00:00:01 gnome-keyring-d
1552 ?      00:00:03 csd-media-keys
1553 ?      00:00:00 at-spi-bus-laun

```

- The result is confusing because you're inside another program now.
- Open a shell (in Emacs with M-x `shell` or a terminal) and type `ps`. You should see something like this:

```

      PID TTY          TIME CMD
12254 pts/1    00:00:00 bash
12257 pts/1    00:00:00 ps

```

- **What this means:**
  - You see two PID - the shell program and the `ps` program
  - TTY ("teletype") is the *controlling terminal* for the process
  - TIME is the amount of CPU time consumed by the process
- Run `ps` again, this time add the option `x`

```
ps x | head -n 10
```

```

      PID TTY          STAT TIME  COMMAND
1313 ?      Ss      0:00  /lib/systemd/systemd --user
1314 ?      S        0:00  (sd-pam)
1323 ?      S<s1    0:00  /usr/bin/pipewire
1324 ?      S<s1    47:15  /usr/bin/pulseaudio --daemonize=no --log-target=journal
1325 ?      Ss1     0:01  cinnamon-session --session cinnamon
1343 ?      Ss      0:04  /usr/bin/dbus-daemon --session --address=systemd: --nofd
1541 ?      SL1     0:01  /usr/bin/gnome-keyring-daemon --start --components=pkcs1
1552 ?      Sl      0:03  csd-media-keys
1553 ?      Sl      0:00  /usr/libexec/at-spi-bus-launcher --launch-immediately

```

- `ps x` (no dash!) shows all processes regardless of what terminal they are controlled by. `?` indicates no terminal (like daemons).
- How many processes that you own that have no terminal?

```
ps x | grep ? | wc -l
```

112

- List only the first 5 lines of the `ps x` listing.

```
ps x | head -5
```

PID	TTY	STAT	TIME	COMMAND
1313	?	Ss	0:00	/lib/systemd/systemd --user
1314	?	S	0:00	(sd-pam)
1323	?	S<sl	0:00	/usr/bin/pipewire
1324	?	S<sl	47:15	/usr/bin/pulseaudio --daemonize=no --log-target=journal

- The column `STAT` reveals the current status of the process, see table .

STATE	MEANING
R	Running or ready to run
S	Sleeping, waiting for an event (e.g. keystroke)
D	Uninterruptible sleep, waiting for I/O (e.g. disk)
T	Stopped, received instruction to stop
Z	Zombie child process, abandoned by parent
<	High priority (not <i>nice</i> - more CPU time)
N	Low priority ( <i>nice</i> ) - served once < are done

There may be more characters denoting exotic process characteristics (see `ps(1)`). E.g. `s` is a *session leader*, `+` is a *foreground* process, and `l` is multi-threaded.

- Check if you have any running processes (`R`) or Zombie processes (`Z`):

```
ps x | grep -cE [RZ]
```

6

- You get even more information with the option `aux`. Redirect the output of `ps aux` to a file `psaux.txt`, and print only the first 5 lines.

```
ps aux | tee psaux.txt | head -5
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.1	166500	8892	?	Ss	Apr19	0:05	/sbin/init splash
root	2	0.0	0.0	0	0	?	S	Apr19	0:00	[kthreadd]
root	3	0.0	0.0	0	0	?	I<	Apr19	0:00	[rcu_gp]
root	4	0.0	0.0	0	0	?	I<	Apr19	0:00	[rcu_par_gp]

- You should see PID 1, the `init` program. The `splash` options means that you can see a splash screen during boot.
- Table shows some header definitions

HEADER	MEANING
USER	User ID - this is the process owner
%CPU	CPU usage in percent
%MEM	Memory usage in percent
VSZ	Virtual memory size (kB)
RSS	Resident set size - RAM use in kB
START	Process starting time and date

- Why is the CPU usage of `init` zero, while the Memory usage is non-zero? How much RAM does the program actually use?

ANSWER: The `init` program only runs during the booting process, but as part of the **kernel** it is loaded into the central memory. It occupies 8MB.

## Viewing processes dynamically

- `ps` provides a snapshot, but `top` provides a real-time view.
- Open a terminal (in or outside of Emacs) and run `top`. You can stop the command with `C-c` or `q`.
- `top` refreshes every three seconds and shows the top system processes. It includes a summary at the top and a table sorted by CPU activity at the bottom.
- The system summary contains a lot of good stuff. Table gives a rundown.
- `top` accepts some keyboard commands like `h` (help) and `q` (quit).

```

top - 21:52:54 up 2 days, 9:47, 1 user, load average: 0.19, 0
Tasks: 180 total, 1 running, 179 sleeping, 0 stopped, 0 zo
%Cpu(s): 0.4 us, 1.7 sy, 0.0 ni, 97.8 id, 0.1 wa, 0.0 hi,
MiB Mem : 3787.2 total, 1964.6 free, 474.2 used, 1348.4
MiB Swap: 100.0 total, 100.0 free, 0.0 used. 2971.7

```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM
1161	pi	20	0	302736	97504	77316	S	5.6	2.5
12974	pi	20	0	11356	3052	2600	R	1.0	0.1
10766	pi	20	0	121660	32344	27112	S	0.7	0.8
12	root	20	0	0	0	0	I	0.3	0.0
321	avahi	20	0	6916	2780	2520	S	0.3	0.1
1258	pi	20	0	291240	73884	58144	S	0.3	1.9
12783	root	20	0	0	0	0	I	0.3	0.0
12926	root	20	0	0	0	0	I	0.3	0.0
1	root	20	0	33832	8784	6868	S	0.0	0.2
2	root	20	0	0	0	0	S	0.0	0.0
3	root	0	-20	0	0	0	I	0.0	0.0
4	root	0	-20	0	0	0	I	0.0	0.0
8	root	0	-20	0	0	0	I	0.0	0.0
9	root	20	0	0	0	0	S	0.0	0.0
10	root	20	0	0	0	0	S	0.0	0.0
11	root	20	0	0	0	0	S	0.0	0.0

Figure 2: Top view

ROW	FIELD	MEANING
1	top	Program name
	21:52:54	Current time of day
	up 2 days 9:49	<i>uptime</i> since last boot
	1 user	No. of users logged in
	load average	No. of processes waiting to run
		Values < 1.0 means not busy
2	Tasks:	No. of processes and their states
		total, running, sleeping, stopped
3	Cpu(s):	Activities that the CPU performs:
		us: user processes (not kernel)
		sy: system processes (kernel)
		ni: nice (low prio) processes
		id: idle processes
		wa: waiting for I/O
4	Mem:	Physical RAM used
5	Swap:	Swap space (virtual memory) used



- `top` is better than any graphical application (e.g. the Task Manager that you have on your Pi) - it is faster and consumes far less resources.

## Controlling processes

### Interrupting a process

- As a guinea pig program, we use `emacs`.
- Open a terminal (inside Emacs after splitting the screen with `C-x 2` or outside of Emacs), and enter `emacs` at the prompt. A new Emacs editor window appears. Notice that the terminal prompt does not return.
- Close the new Emacs editor manually by clicking on the **X** in the upper right corner. The prompt in the Shell returns.
- Enter `emacs` again in the shell, and interrupt it with `CTRL-C` (outside of Emacs, or with `C-c C-c` on the Emacs `*shell*`).
- Many programs can be interrupted this way by sending an **interrupt** signal to the **kernel**.

### Putting a process in the background

- The terminal has a *foreground* and a *background*. To launch a program so that it is immediately placed into the background, follow it with an ampersand `&` character
- Start Emacs from the shell in the background. An Emacs window should open. Look at the terminal.
- The message that appeared is part of shell *job control*. It means that we have started job number 1 with the PID 13899. If you check the process table with `ps`, you should see the process

```
[1] 13899
```

- `grep` the `emacs` process from the process table using the PID.

```
13928 pts/1    00:00:04 emacs
```

- The `jobs` command lists the jobs that were launched from our terminal. Try it. You should see something like this:

```
[1]+  Running                  emacs &
```

## Returning a process to the foreground

- A process in the background is immune from keyboard input - you cannot interrupt it with `CTRL-C`. To return it to the foreground, use the `fg` command.
- On the shell where you started it, return the process to the foreground with the command `fg %1`. The `1` is the `jobspec`.
- Kill the Emacs process with `C-c C-c` or `CTRL-C` on the shell where you started it.
- If you enter `jobs` you get no response, and `fg` tells you there's no job.

## Stopping or pausing a process

- Start an `emacs` process in a terminal (NOT in an Emacs shell) - it's now in the foreground. If you press `CTRL-z` in the shell, the process is stopped.

```
pi@raspberrypi:~ $ emacs
^Z
[1]+  Stopped                  emacs
pi@raspberrypi:~ $
```

- To bring the process back, you can either bring it into the foreground with `fg %1`, or resume the process in the background with `bg %`. Try both.
- **Why would you launch a graphical program from the shell?**
  - The program may not be listed in the GUI
  - You see error messages that otherwise are invisible
  - Some graphical programs have useful command line options

## Killing a process

- `kill` is used to terminate processes using the PID. Start Emacs from the shell *in the background* (inside or outside of emacs), and then kill it with `kill PID`.

*Tip: you get the PID with `ps`, or right after executing the background command.*

- `kill` does actually not "kill" the process, it sends it a signal. We have already used some of these signals:

SIGNAL	MEANING
INT	CTRL-C - interrupt process
TSTP	CTRL-Z - terminal stop
HUP	Hang up (used by daemons)
KILL	Kill without cleanup
TERM	Terminate with <code>kill</code>
STOP	Stop without delay

- Some of these signals are sent to the target program (identified by PID) while others are sent straight to the kernel.

## More process commands

Some fun commands to play with and explore. We already looked at `ps`. You may have to install these.

COMMAND	MEANING
<code>ps</code>	Process list arranged as tree pattern
<code>vmstat</code>	System usage snapshot
<code>xload</code>	Draws a graph showing system load over time
<code>tload</code>	Draws graph in terminal

## Summary

- Multitasking by rapidly switching tasks, managed by the kernel.
- Memory layout in processes includes compiled code, initialized/uninitialized data, stack, and heap.

- Useful shell commands for process management: `ps`, `top`, `jobs`, `bg`, `fg`, `kill`, `shutdown`.
- Kernel starts with `init` program to launch system services.
- Services typically run as background daemons, managed via `systemctl`.
- Processes are tracked via Process IDs (PID).
- Snapshot of processes using `ps`, dynamic view with `top`.
- Process statuses include running, sleeping, stopped, and zombie states.
- Interrupt, background, and foreground control of processes with commands like `CTRL-C`, `&`, `fg`, `bg`.
- `kill` command for sending signals to processes.
- Additional tools: `vmstat`, `xload`, `tload` for system performance analysis.
- Offers enhanced control and visibility, crucial for system optimization and troubleshooting.

## References

- Silberschatz, Galvin and Gagne (2018). Operating System Concepts - 10th edition, Wiley.
- Shotts, The Linux Command line (2019). NoStarch.